DB2 Data Management Software





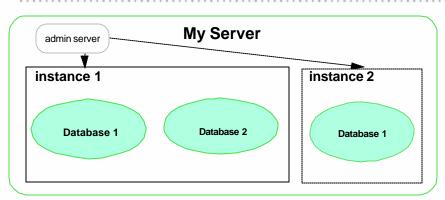
DB2 Version 8 Basic Architecture

DB2 Quickstart Education Maintained by Paul Yip (ypaul@ca.ibm.com)

February 2003

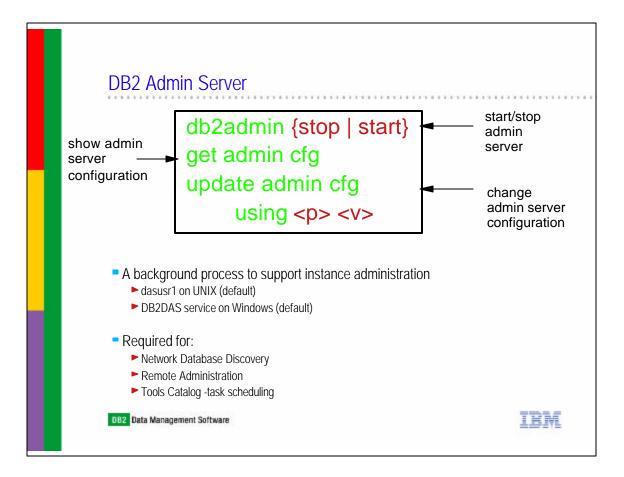
IBM Software Group

The DB2 Server Picture



- One physical machine has:
 - One Admin Server
 - One or multiple instances
- Each instance (database manager)
 - is a logical server context
 - can have one or many <u>databases</u>

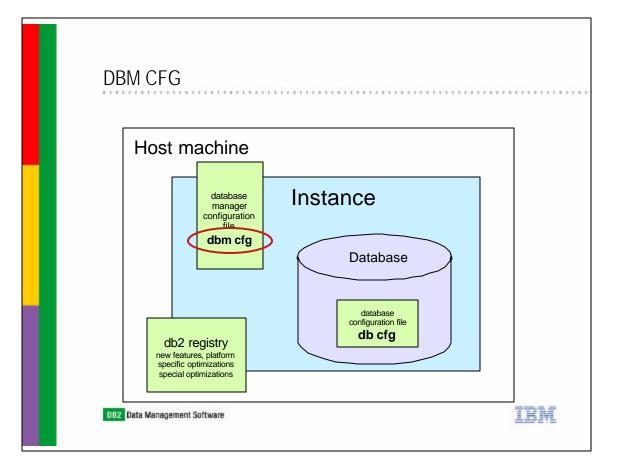
DB2 Data Management Software

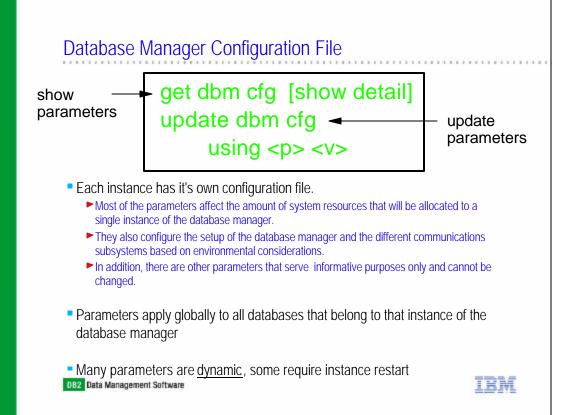


Instance Concepts

- An instance, also known as <u>Database Manager</u>
 - ► Manages the system resources assigned to it
 - ► Can be independently stopped or started
 - ► has its own set of processes
 - ► Has its own configuration file, <u>DBM CFG</u>.
 - security plan
 - performance
 - communications
- example: One instance for development, another for QA testing. Fully Independent.

DB2 Data Management Software





Common Instance (DBM CFG) Parameters (cont..)

DIAGLEVEL

► Diagnostic log (db2diag.log) detail level

NOTIFY LEVEL

► Notification log (*.nfy) detail level

MON_HEAP_SZ

- ► Database performance monitor heap size
- ▶ Default generally too small recommend 128 or higher

■ INTRA PARALLEL

- ► Enable/disable intra-partition parallelism
- ► Keep disabled for OLTP workloads

■ SVCENAME

- ► TCP/IP Service Name, as mapped in /etc/services file
- ► Must be set to support remote communications





Common Instance (DBM CFG) Parameters

MAXAGENTS

► Maximum number of database manager agents available at any given time to accept application requests. This value limits the memory usage of the instance, because each additional agent requires additional memory.

NUM_POOLAGENTS

- ► A guideline for how large the agent pool should grow. The pool can save the cost associated with allocating an agent.
- ► Not significant in environments where connection pools are used

NUMDB

- Limits the maximum number of concurrently active databases.
- ► Value should reflect reality as closely as possible.

■ SHEAPTHRES

- ► Controls the total amount of memory for sorting available across the entire instance for all sorts
- ► start with roughly 10x SORTHEAP (DB CFG)

DB2 Data Management Software

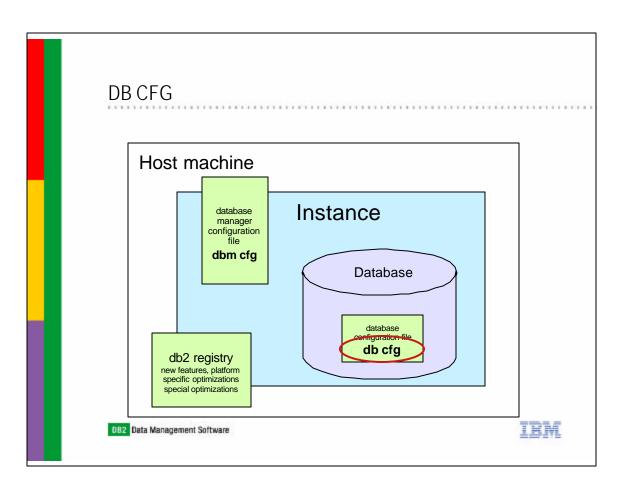


Databases Concepts

- A collection of tables and related objects.
- Each database includes
 - system catalog tables (database meta data)
 - ► a database configuration file (db cfg)
 - ► a recovery log
- Databases are isolated from each other

DB2 Data Management Software





Database Configuration File

show parameters get db cfg [show detail] update db cfg for <dbname> update update parameters

- Parameters that affect the database.
 - if there are two databases in an instance, there will be two database configuration files
- Parameters include:
 - ► location of database logs
 - defaults for query optimization
 - ► size of memory for lock management
 - ►etc..
- Many parameters are dynamic. Others may require all applications disconnect. Changes to take effect on first re-connection





Common Database (DB CFG) Parameters

- SORTHEAP
 - ► Defines the maximum number of memory pages for sorts
 - Set this parameter big enough to avoid sort overflows, but not so big that the optimizer attempts too many sorts
- MAXAPPLS
 - ► Specifies the maximum number of applications (local and remote) that can be connected to the database
- AVGAPPLS
 - ► Specifies the expected number of concurrently executing (large query) applications

 affects how much buffer pool the optimizer expects to have available
- NEWLOGPATH
 - ► change the location of the db2 logs.
 - ► recommendation: always move logs onto separate, dedicated disks
- MIRRORLOGPATH
- DB2 upath to mirrored logs



Common Database (DB CFG) Parameters (cont...)

CHNGPGS_THRESH

- ► Specifies the percentage of the buffer pool which must be "dirty" before page cleaners are triggered
- ► frequent, small writes vs. infrequent, large writes

■ NUM IOCLEANERS

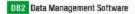
- ► Specifies the number of asynchronous page cleaners
- Asynchronous page cleaners can save an agent from having to write changed data to the disk when it needs to bring data into the buffer pool

NUM_IOSERVERS

- ► Specifies the number of I/O servers for a database
- ► I/O servers perform prefetching which can reduce the amount of I/O that agent has to perform to execute a query

LOGBUFSZ

- ► Specify buffer size for database logging
- ► defaults generally too small





Common Database (DB CFG) Parameters (cont...)

UTIL_HEAP_SZ

►usually increased when tuning for utilities such as LOAD

LOCKLIST

► Memory storage size for lock list

MAXLOCKS

► Maximum percent of LOCKLIST that can be used by an application

■ SHEAPTHRES SHR

► Sortheap threshold for shared sorts (for intra-partition parallelism)

SORTHEAP

► maximum size of sort heap that a db2 agent can allocate

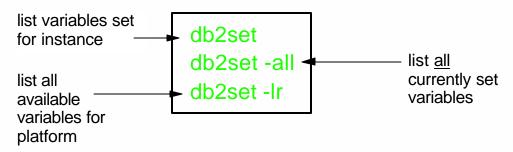
LOCKTIMEOUT

- maximum time an application will wait on locks before rolling back
- ► Default is wait forever. Therefore, should be changed to reflect needs

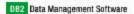
 □B2 □B12 □B13 Management Software







- A place to specify environment or platform specific settings
 - ► e.g. settings for RAID, optimizer paths not yet fully integrated, communication protocols, environment specific performance enhancements
- Settings can be applied to various levels
 - **►** environment
 - ►instance
 - ► global (all instances)





Database Objects

- Table
 - ► A collection of data logically arranged in columns and rows
- View
 - ► A representation of the data in tables
- Index
 - ► An ordered set of keys each pointing to a row in a table
- Schema
 - ► All DB2 UDB objects have a two part name and the schema is the first half of that name.
 - ► A schema is an identifier that helps group database objects together.
 - ► A schema can be owned by an individual who can then control access to the objects within it.
 - ► Schemas can be either implicit or explicit when accessing an object.
- Database Application Objects
 - ► Sequences
 - ▶ Triggers
 - ► User Defined Functions (UDFs)
 - ► Stored Procedures

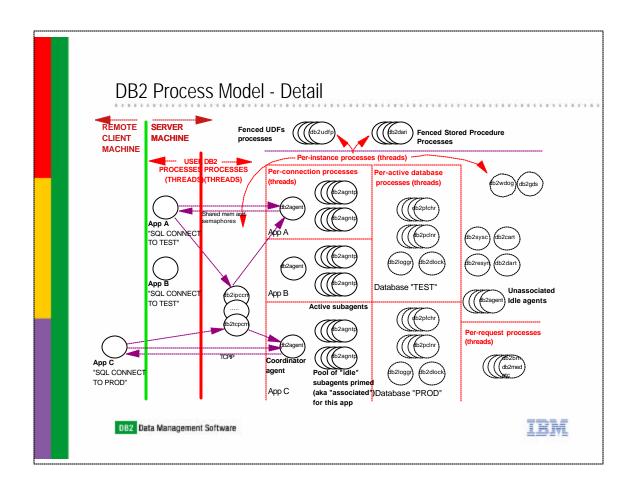
DB2 Data Management Software

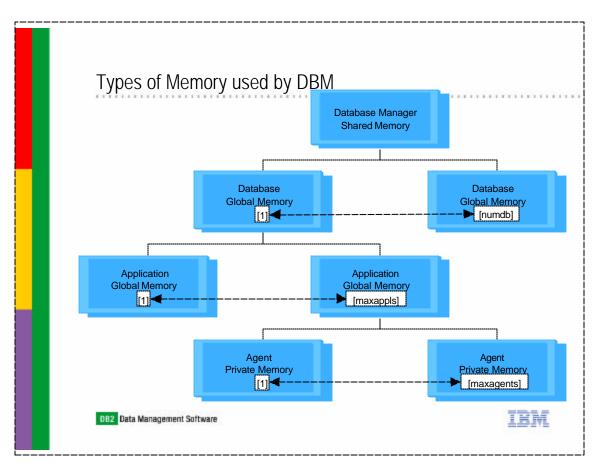
DB2 Object Heirarchy dbm configuration file Instance Catalog Catalog <u>Databas</u> <u> Database</u> 2 Lock List Logs DB config file DB config file Lock List Table1 Table2 Table2 Table3 Index1 View1 View1 Index1 **BLOBs** Table View3 spaces IBM

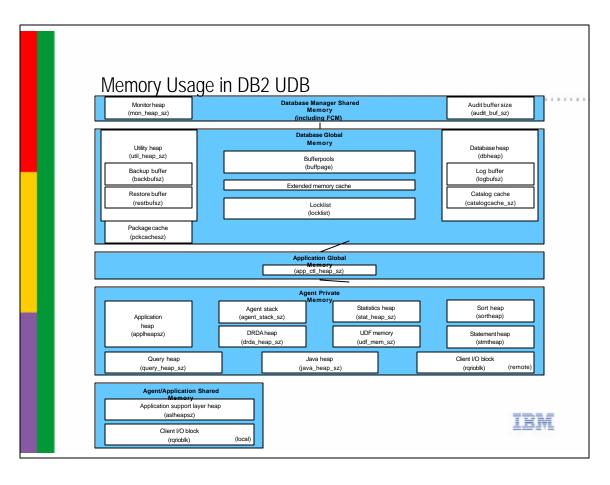
Common Processes

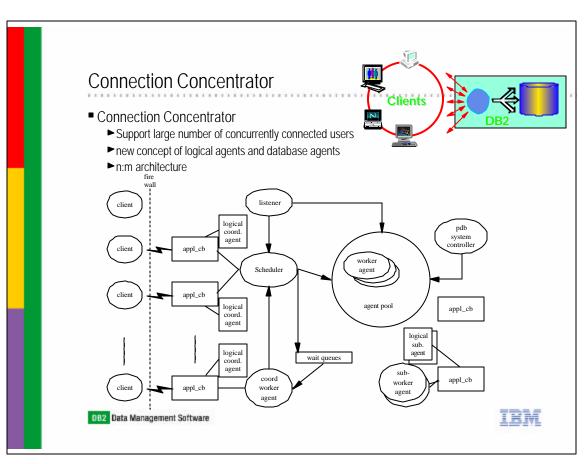
- db2sysc
 - ► database instance process (one per INSTANCE)
- db2tcpcm
 - ► tcpip communication/listener
- db2agent
 - ► Coordinator agent that performs database operations on behalf of applications (at least 1 per connection, depending if Connection Concentrator is enabled).
- db2agntp
 - Active subagent spawned if INTRA_PARALLEL is set to YES. Will perform database operations for the application. db2agent will coordinate the work between the different db2agntp subagents.
- db2pfchr
 - ► DB2 asynchronous I/O data prefetcher (NUM_IO_SERVER)
- db2pclnr
 - ► DB2 asynchronous I/O data writer (NUM_IOCLEANER)

DB2 Data Management Software









Database Agents

- Communication between the database and client
 - ► Unix processes vs. Intel threads
- Connection Concentrator is enabled when: max_connections > max_coordagents (DBM CFG)
- Processes/threads
 - ► Logical agent
 - Represents a connected application to the database manager
 - Contains information and control blocks required by an application
 - Controlled by MAX_COORDAGENTS (DBM CFG)
 - ► Worker agent
 - Carries out application requests
 - No permanent attachment to any particular application





Query Parallelism

- Inter-query Parallelism
 - Multiple queries running at the same time
- Intra-query Parallelism
 - Intra-partition parallelism
 - Inter-partition parallelism



